



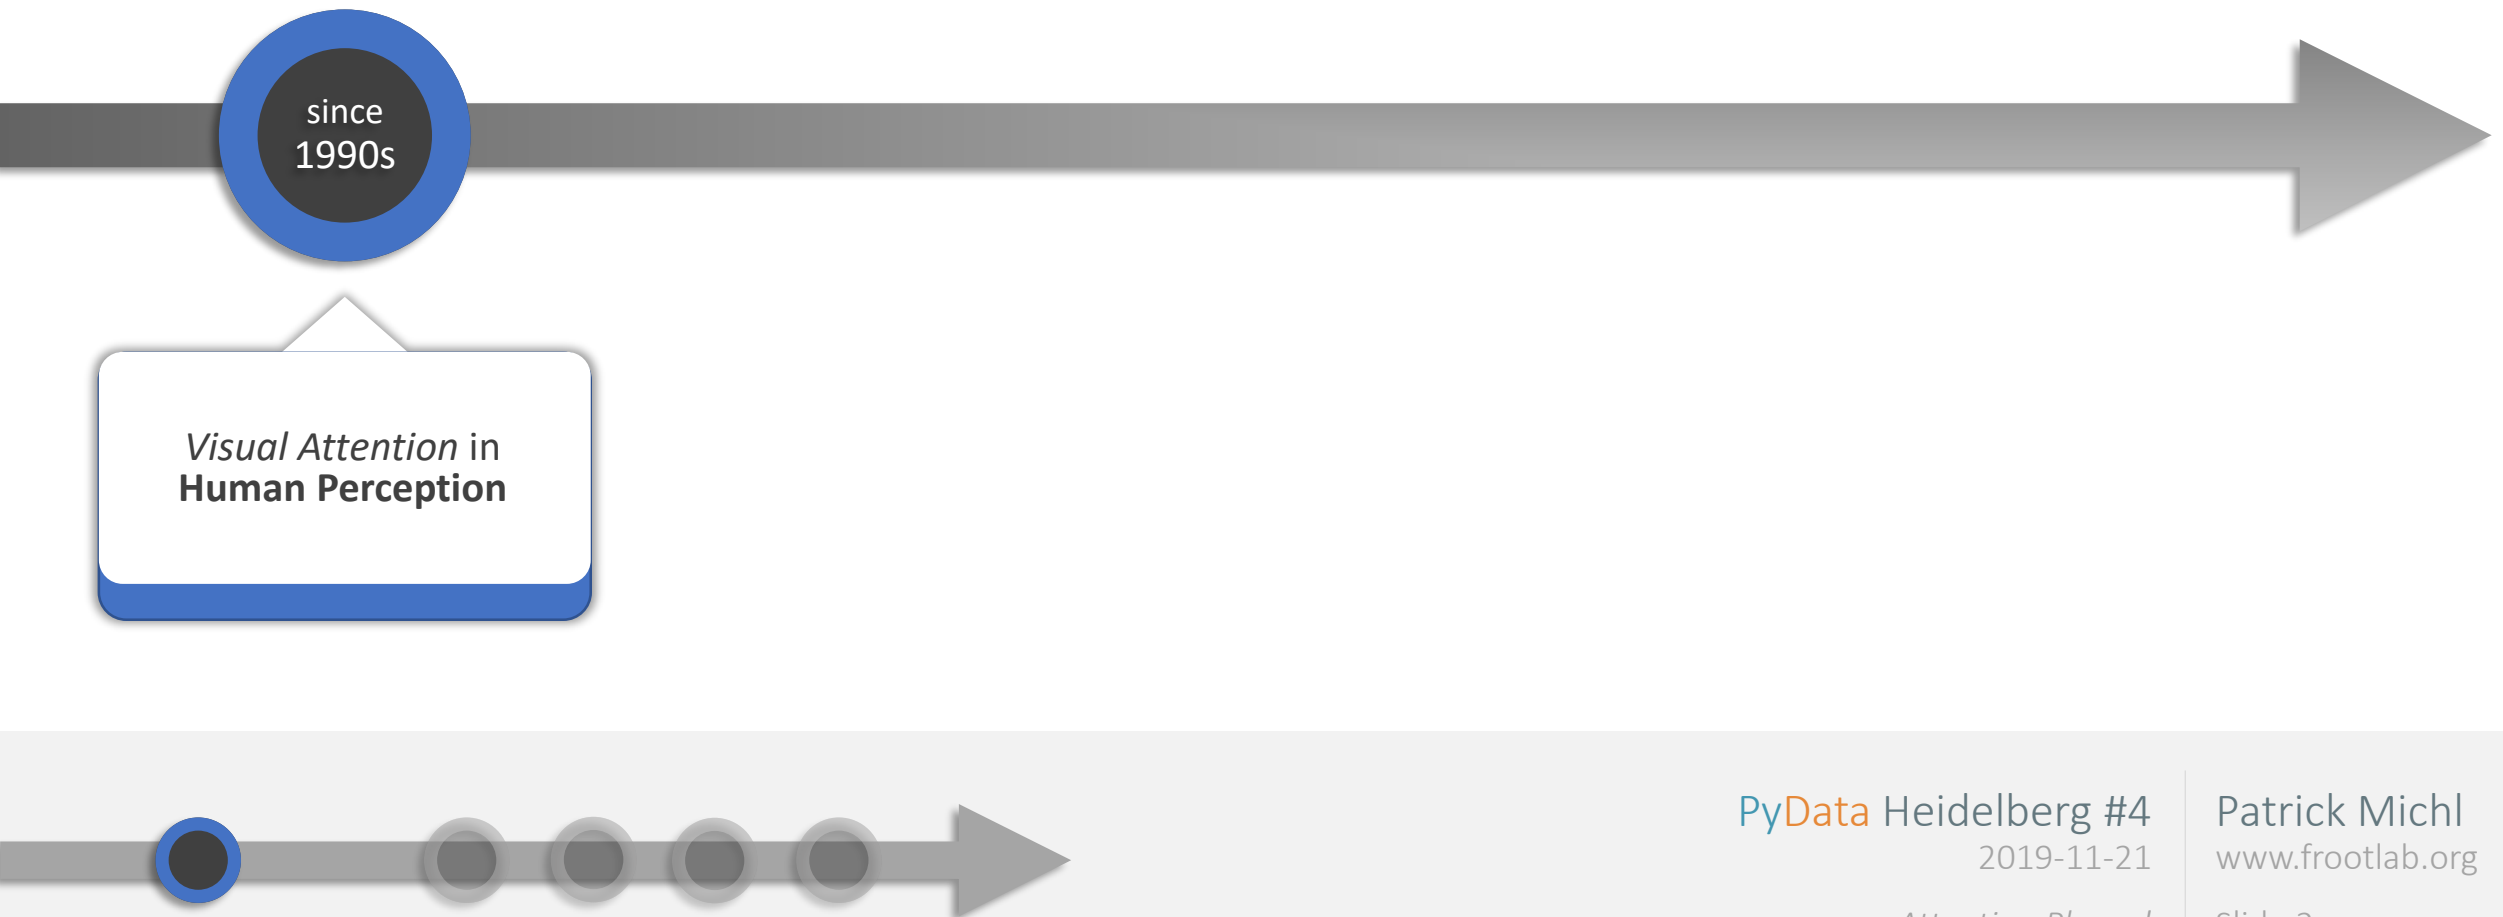
Attention Please!

Patrick Michl
www.frootlab.org

Attention Mechanisms
in Neural Networks

PyData Heidelberg #4
2019-11-21

Milestones in Attention Research



Visual Attention in Human Perception

Visual attention in human perception is based on the dynamics between **Recognition** and **Selection**



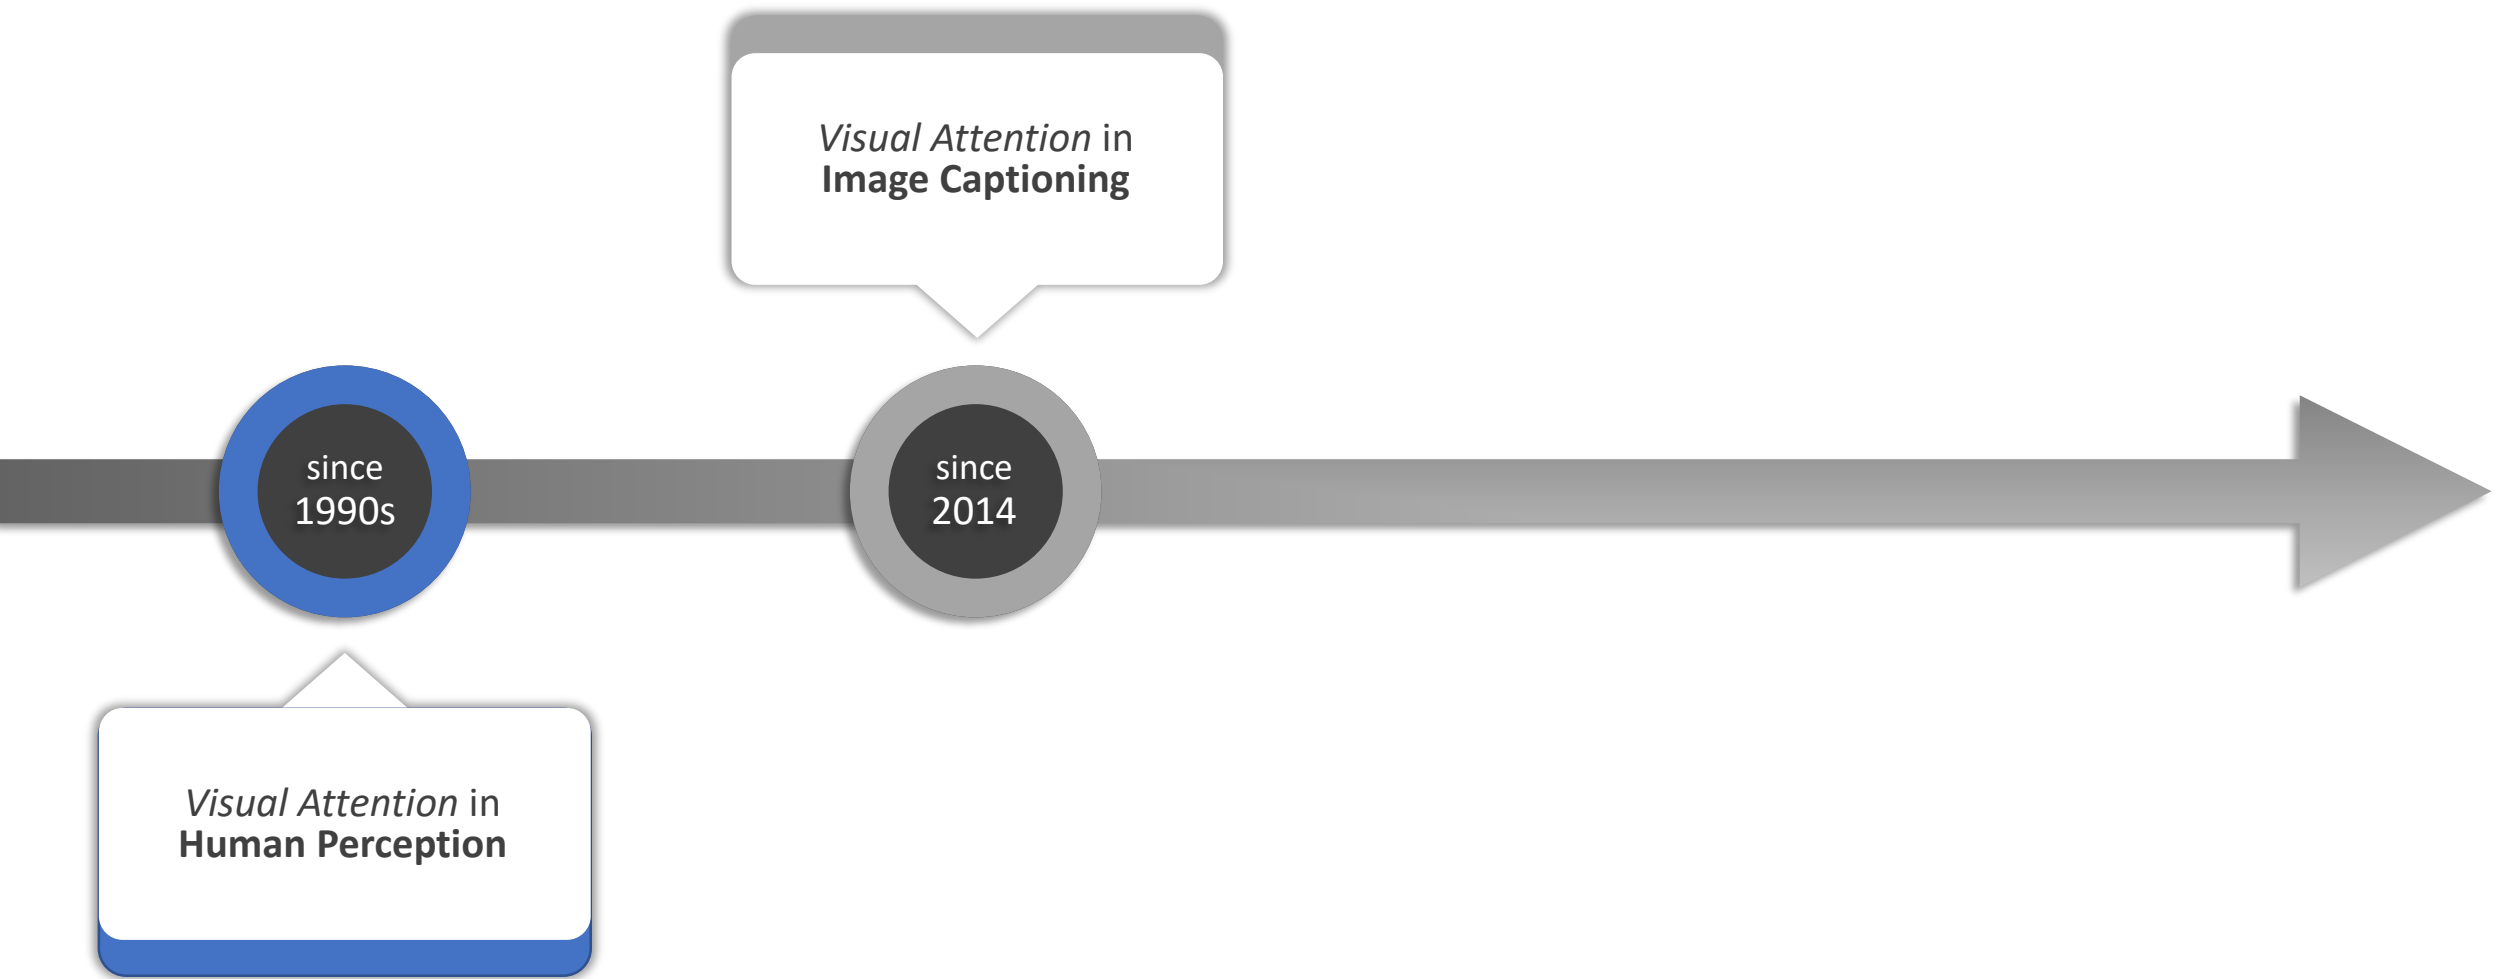


Image source: *John Henderson and Taylor Hayes, UC Davis*



Visual Attention in Image Captioning


Encoder: CNN

Decoder: RNN


Architecture



Encoder



Decoder



A bird on a branch


Visual Attention in Image Captioning

Encoder: CNN

Convolutional features
entangle spatial with
semantic information

Decoder: RNN

Encoder



Visual Attention in Image Captioning

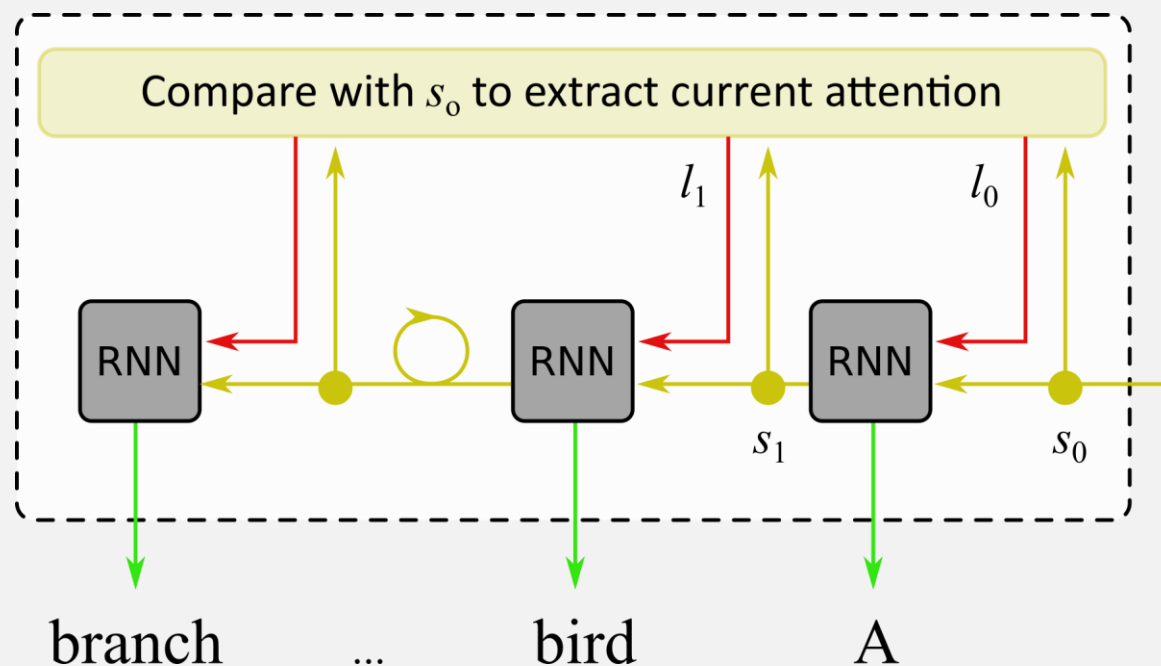
Encoder: CNN

Convolutional features
entangle spatial with
semantic information

Decoder: RNN

Visual Attention
allows to recover spatial
relations of current outputs

Decoder







Visual Attention in Image Captioning


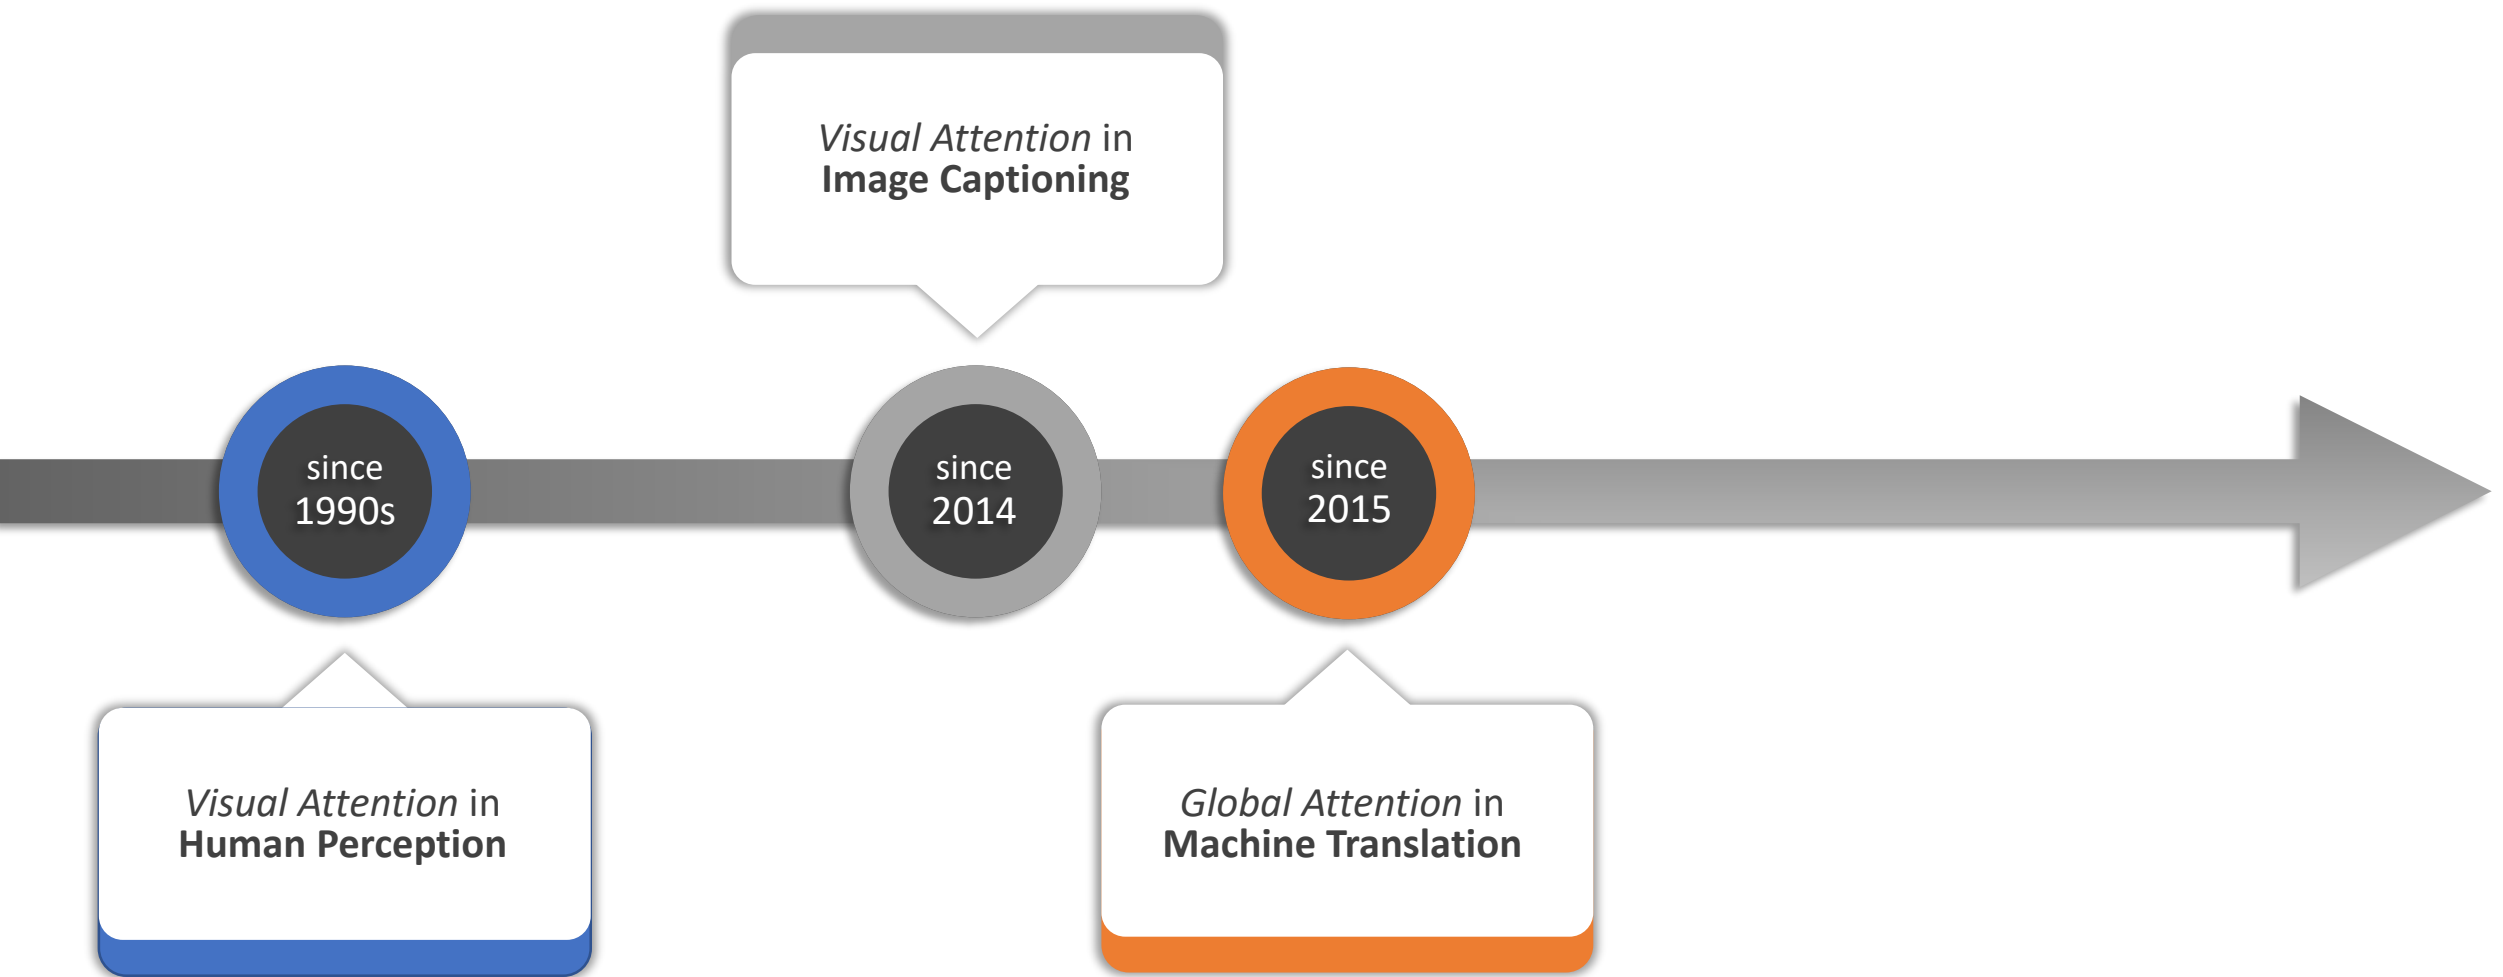
Encoder: CNN

Convolutional features
entangle spatial with
semantic information

Decoder: RNN

Visual Attention
allows to recover spatial
relations of current outputs



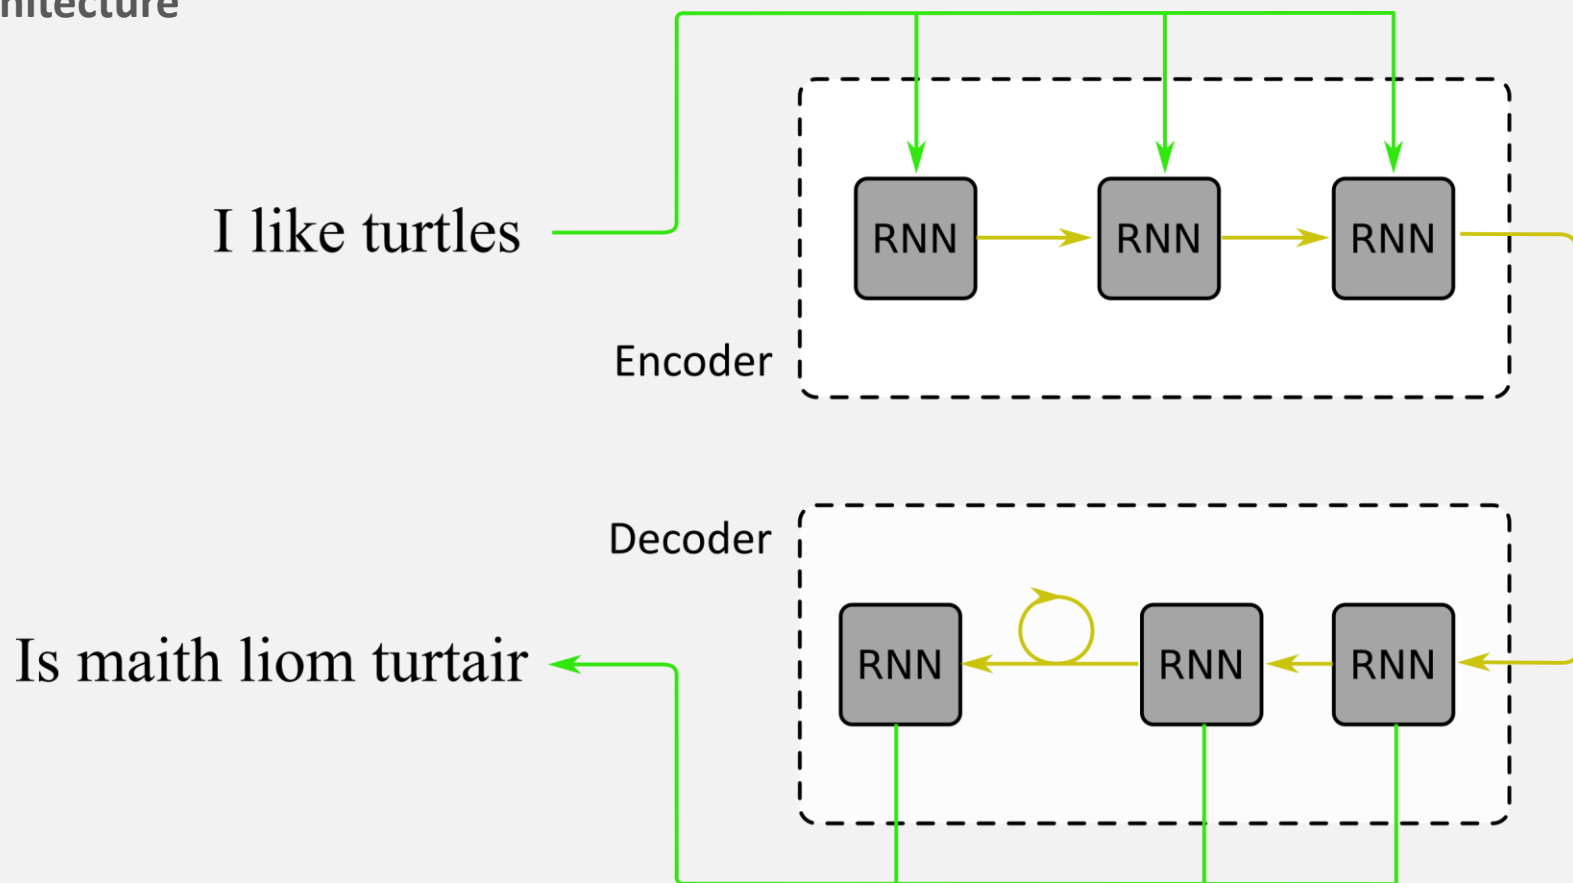


Global Attention in Machine Translation

Encoder: RNN

Decoder: RNN

Architecture



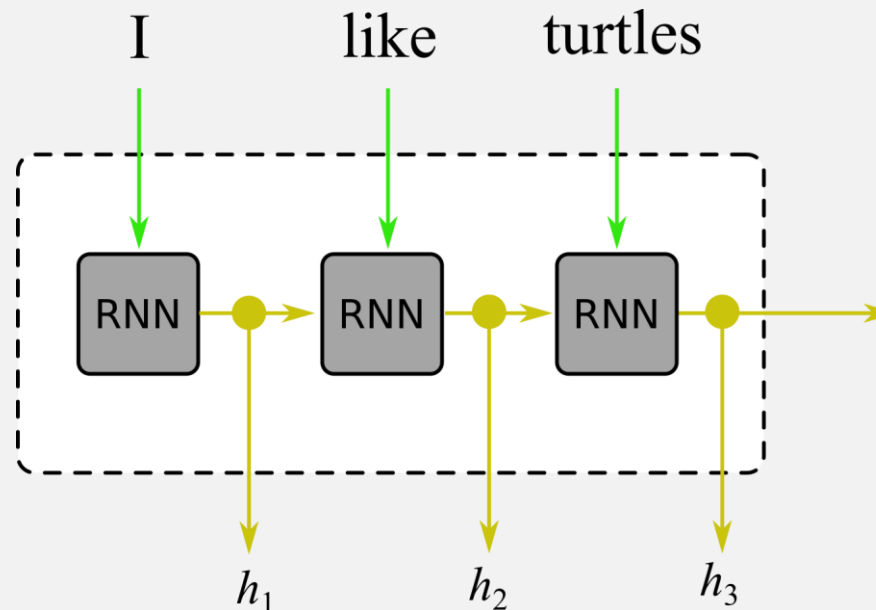
Global Attention in Machine Translation

Encoder: RNN

Sequence of hidden states
encodes the context building
stack

Decoder: RNN

Encoder



Global Attention in Machine Translation


Encoder: RNN

Sequence of hidden states encodes the context building stack

Decoder: RNN

Global Attention allows to recover the context of current outputs

Decoder



Global Attention in Machine Translation

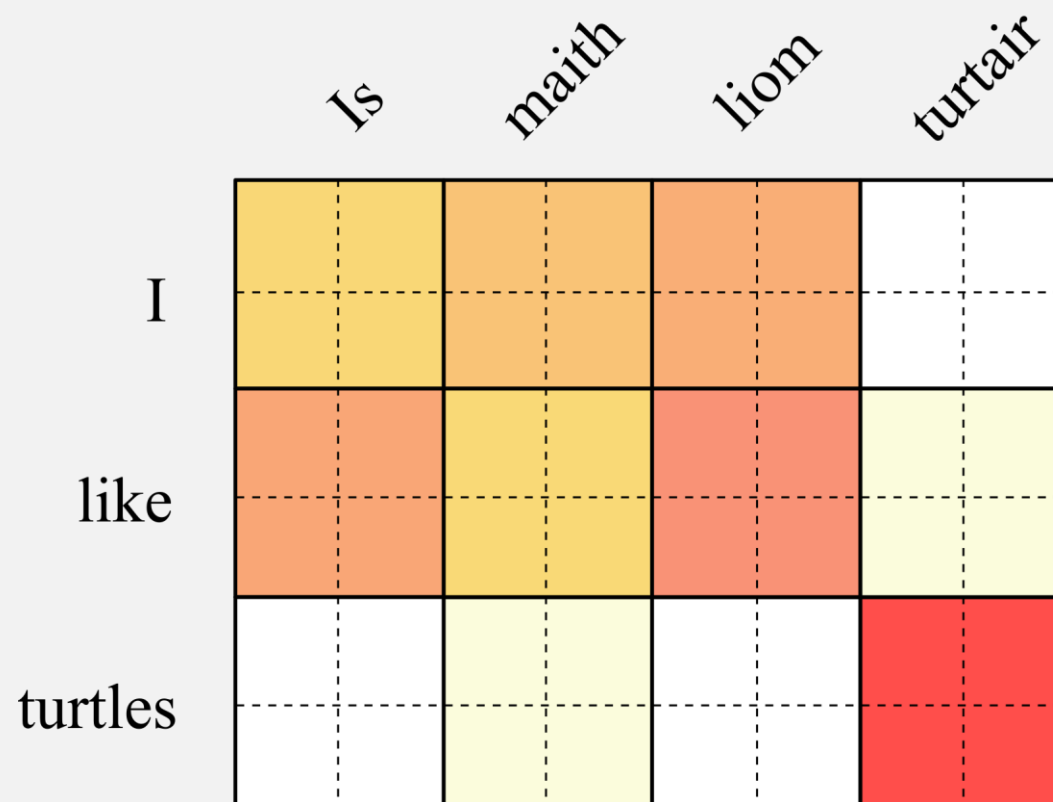
Encoder: RNN

Sequence of hidden states
encodes the context building stack

Decoder: RNN


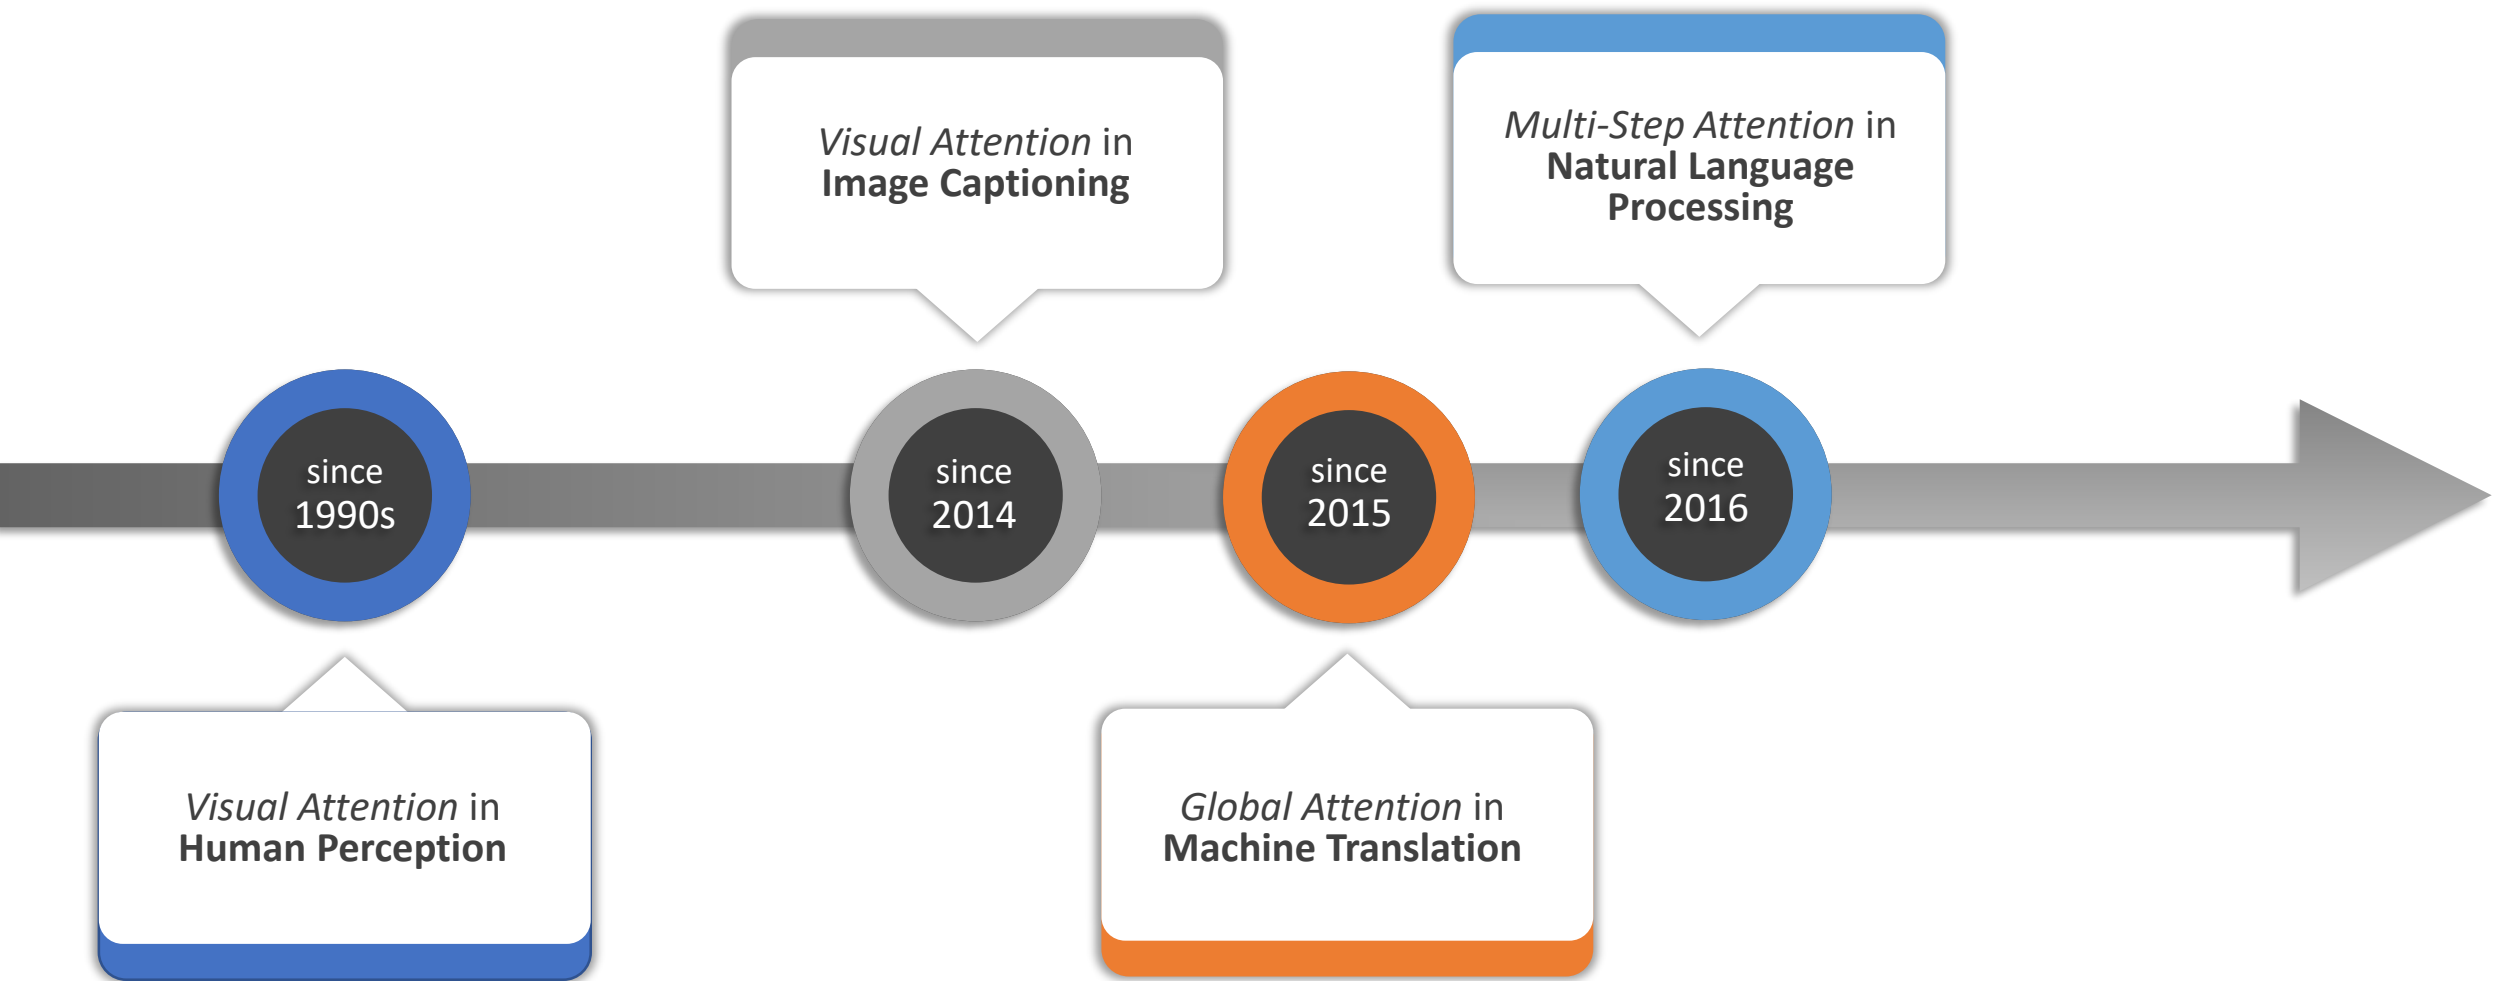
Global Attention
allows to recover the context of current outputs

Context



Important

Unimportant



Multi-Step Attention in Natural Language Processing


Encoder: CNN

Decoder: CNN


Architecture

I like turtles,
so what do I like?

Encoder



Decoder



You like turtles




Multi-Step Attention in Natural Language Processing

Encoder: CNN

**Sequence of convolutional
features** encodes the context
building stack

Decoder: CNN



Multi-Step Attention in Natural Language Processing


Encoder: CNN


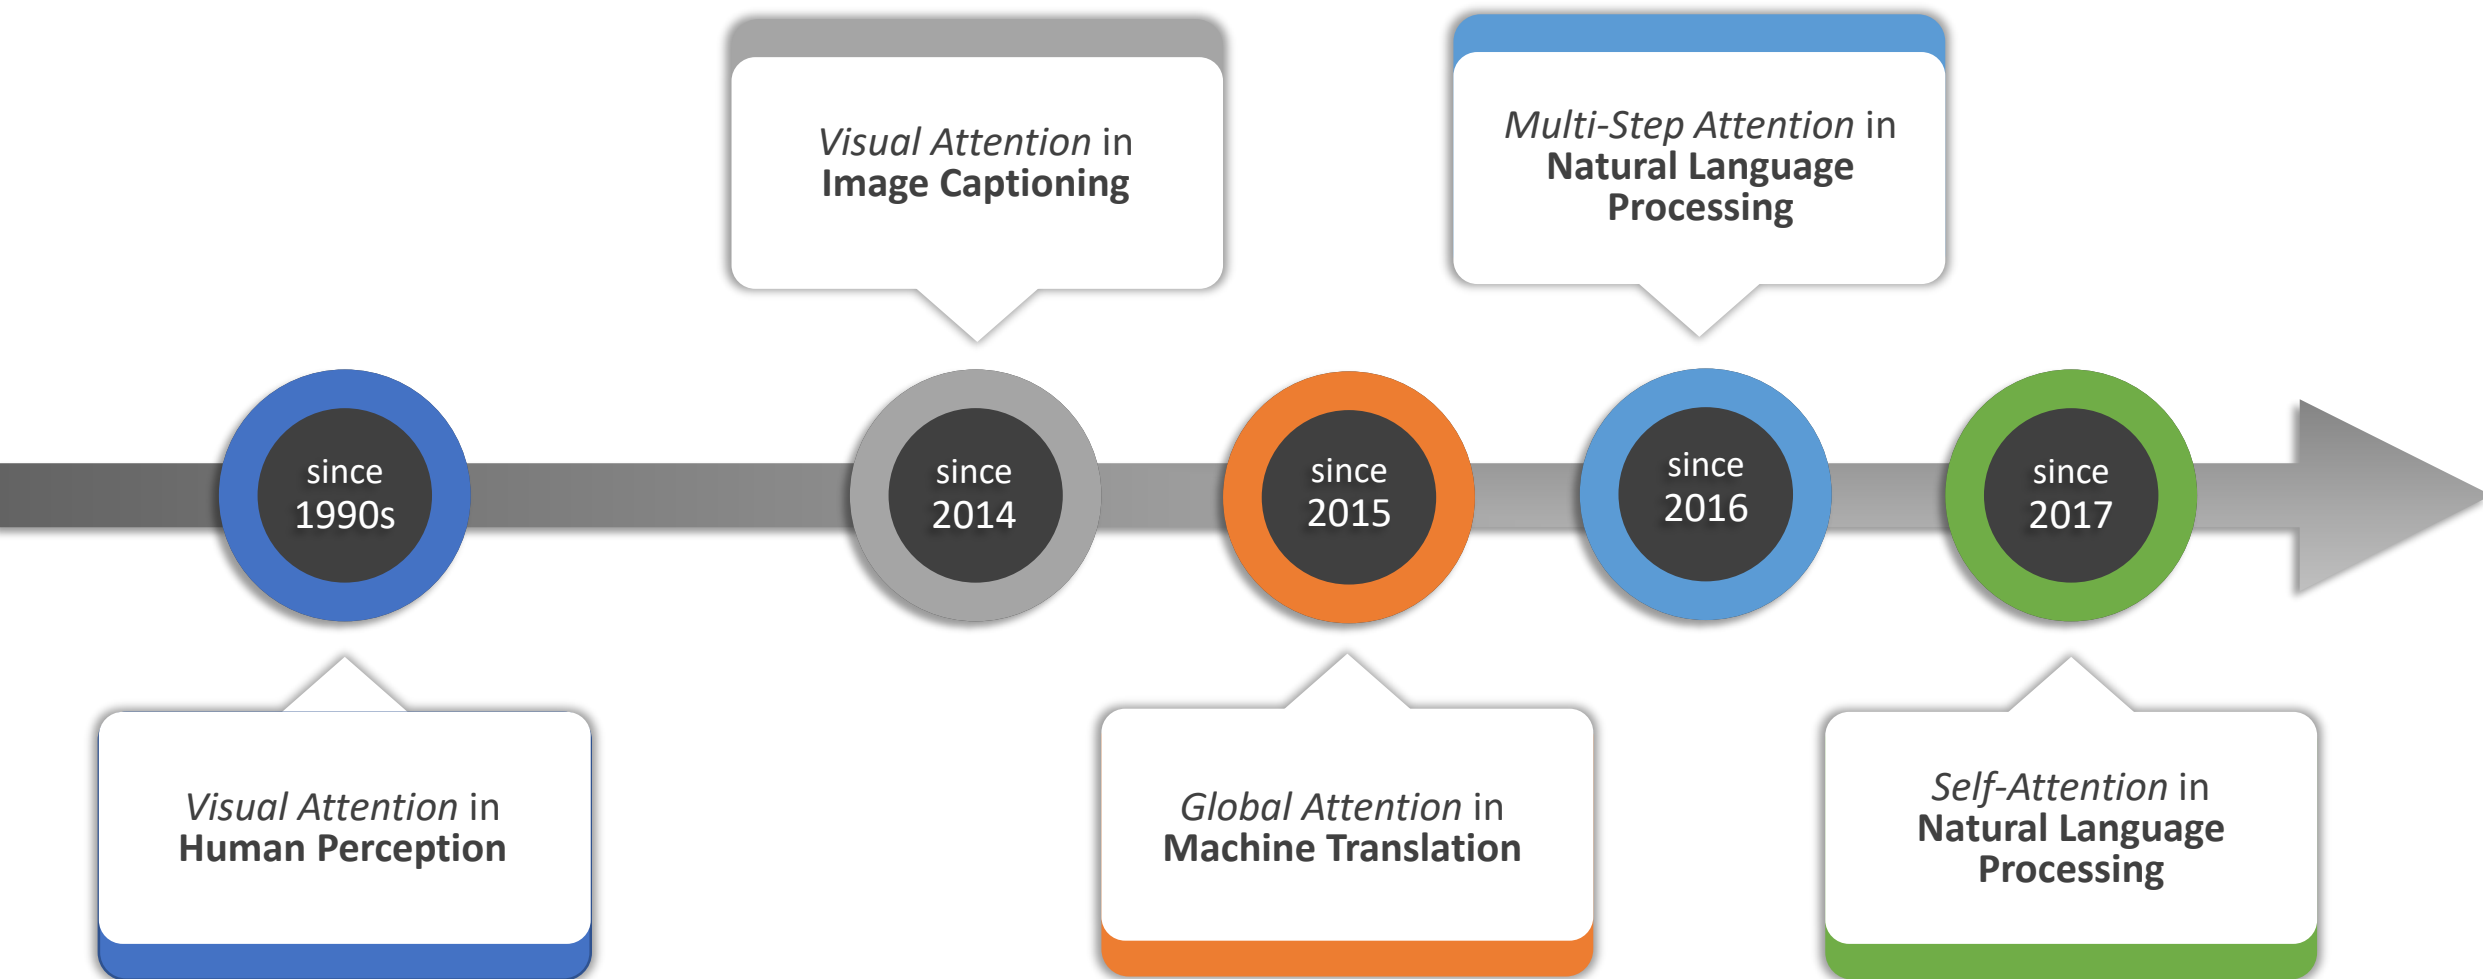
**Sequence of convolutional
features** encodes the context
building stack

Decoder: CNN

Multi-Step Attention
allows to recover the current
context of current
deconvolutional steps

Decoder





Self-Attention in Natural Language Processing


Encoder: Stacked ANN

Decoder: Stacked ANN

Architecture

Cool kids like turtles
and he's a cool kid,
so what does he like?

He likes turtles



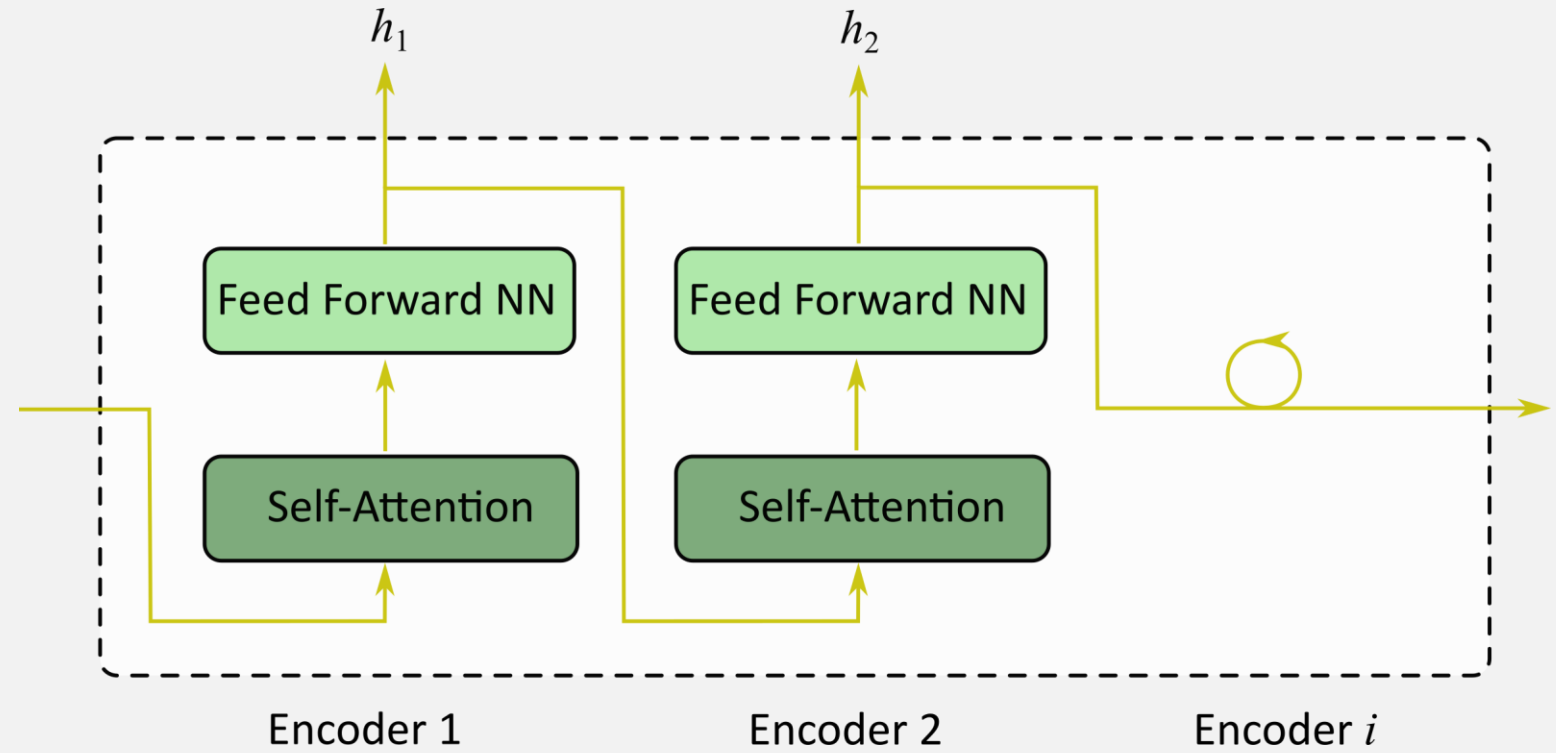
Self-Attention in Natural Language Processing

Encoder: Stacked ANN

Sequence of hidden features
captures the encoder context
building stack

Decoder: Stacked ANN

Encoder Stack



Self-Attention in Natural Language Processing

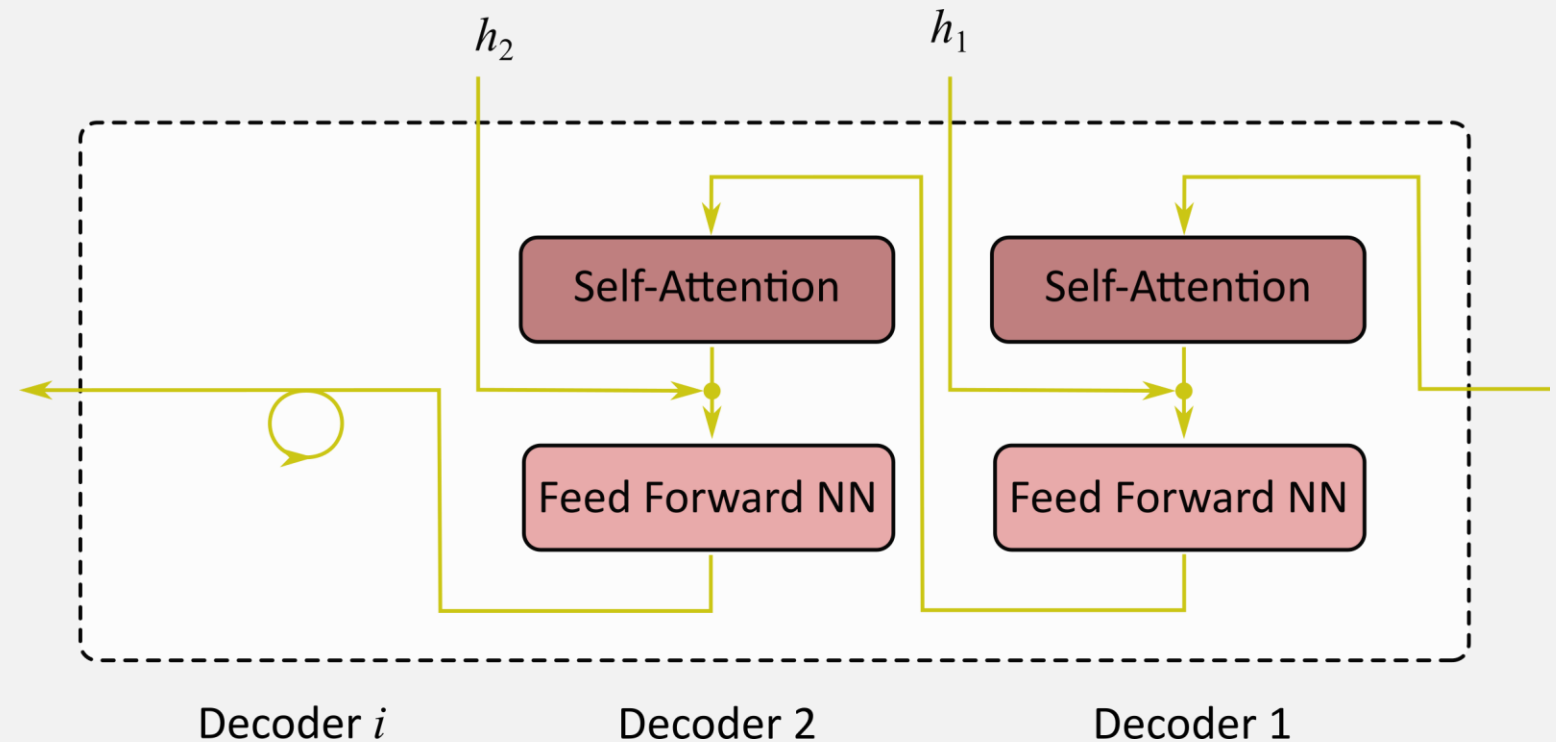
Encoder: Stacked ANN

Sequence of hidden features
captures the encoder context
building stack

Decoder: Stacked ANN

Multi-Step Attention unrolls
the context stack to the
decoder

Decoder Stack



Self-Attention in Natural Language Processing

Encoder: Stacked ANN

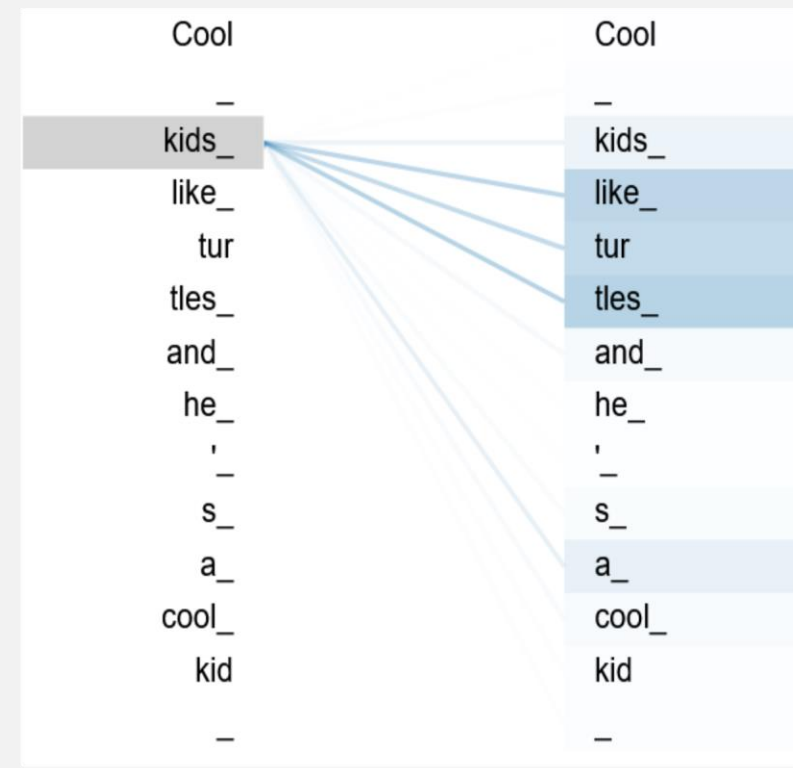
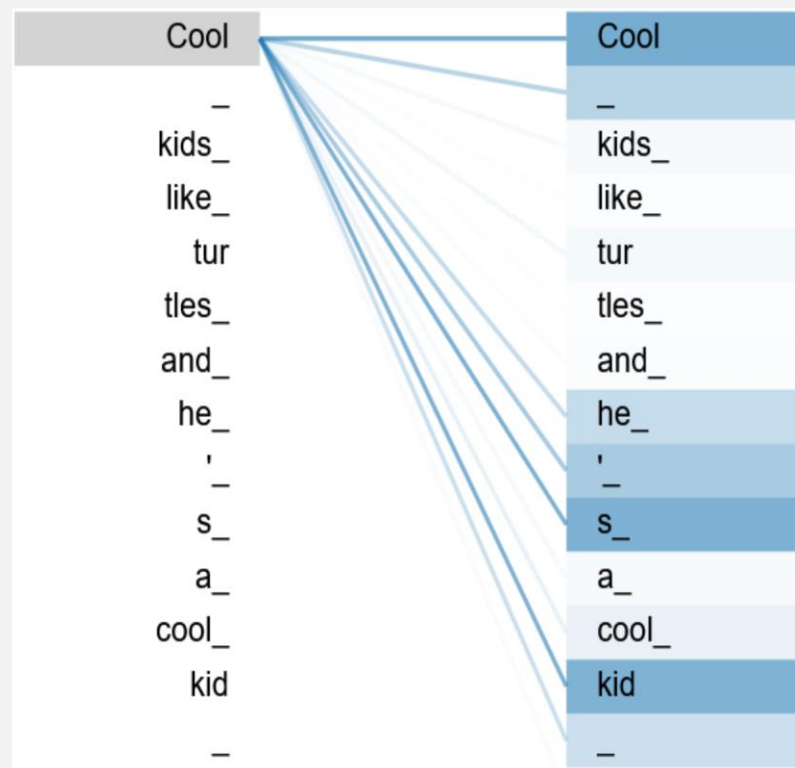
Sequence of hidden features captures the encoder context building stack

Decoder: Stacked ANN

Multi-Step Attention unrolls the context stack to the decoder

Self-attention aggregates context dependencies within the inputs

Self-Attention



Summary

#1

Human Perception is based on the dynamics between *selection* and *recognition*

#2

Attention mechanisms immitate this behaviour by using intermediate states, that entangle *context* with *semantic* information

#3

The incorporation of context provides *dynamic features* that are context specific and therefore improve the model performance

#4

After all - attention mechanisms can also help to *understand the decisions* of deep networks

Summary

#1

Human Perception is based on the dynamics between *selection* and *recognition*

#2

Attention mechanisms immitate this behaviour by using intermediate states, that entangle *context* with *semantic* information

#3

The incorporation of context provides *dynamic features* that are context specific and therefore improve the model performance

#4

After all - attention mechanisms can also help to *understand the decisions* of deep networks

Summary

#1

Human Perception is based on the dynamics between *selection* and *recognition*

#2

Attention mechanisms immitate this behaviour by using intermediate states, that entangle *context* with *semantic* information

#3

The incorporation of context provides **dynamic features** that are context specific and therefore improve the model performance

#4

After all - attention mechanisms can also help to **understand the decisions** of deep networks

Summary

#1

Human Perception is based on the dynamics between *selection* and *recognition*

#2

Attention mechanisms immitate this behaviour by using intermediate states, that entangle *context* with *semantic* information

#3

The incorporation of context provides **dynamic features** that are context specific and therefore improve the model performance

#4

After all - attention mechanisms can also help to **understand the decisions** of deep networks

Thank you for your attention!

