


Structure learning with deep neuronal networks

6th Network Modeling Workshop, 6/6/2013

Patrick Michl




Autoencoders

Biological Model

Validation & Implementation


Dataset

Model




Real world data usually is **high dimensional** ...

Dataset




Model



... which makes **structural analysis** and modeling complicated!

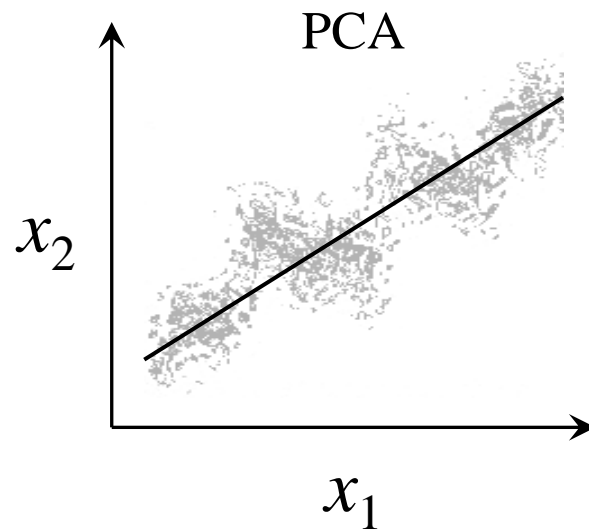
Dataset

Model




Dimensionality reduction techniques like **PCA** ...

Dataset




Model



... can not preserve **complex structures!**


Dataset

Model




Therefore the analysis of **unknown structures** ...

Dataset




Model



... needs more considerate **nonlinear techniques!**


Autoencoder

- Artificial Neuronal Network




Autoencoders are **artificial neuronal networks** ...

Perceptron




Gaussian Units

Gauss Units



input data \mathbf{X}




output data \mathbf{X}'

nal networks ...

Autoencoder


- Artificial Neuronal Network



Autoencoders are **artificial neuronal networks** ...

Autoencoder

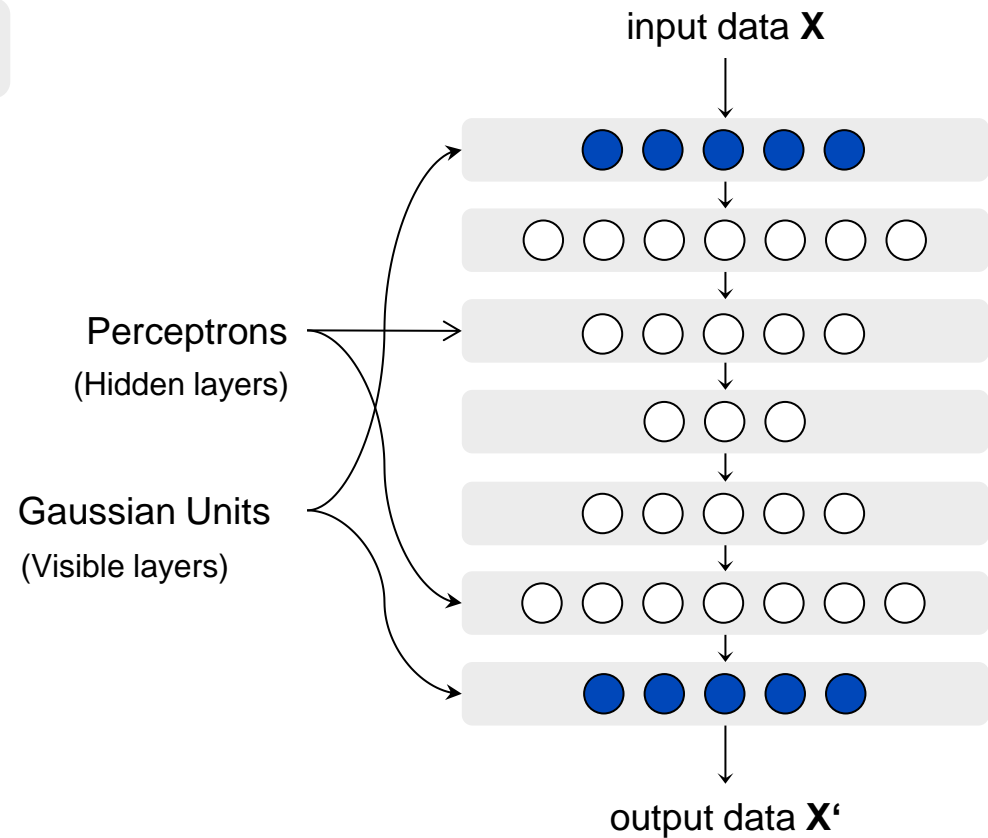
- Artificial Neuronal Network
- Multiple hidden layers



... with **multiple hidden layers**.

Autoencoder

- Artificial Neuronal Network
- Multiple hidden layers



Such networks are called **deep networks**.

Autoencoder


- Artificial Neuronal Network
- Multiple hidden layers

Definition (*deep network*)

Deep networks are artificial neuronal networks with multiple hidden layers

Neurons
(layers)

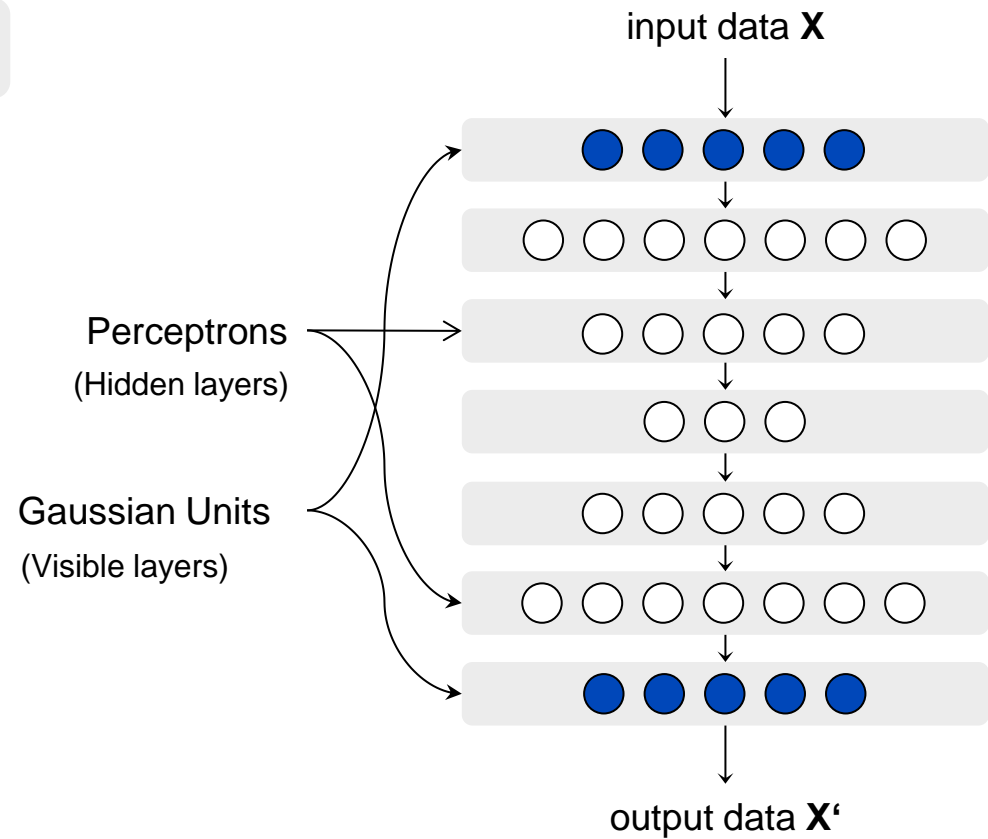
Units
(layers)



Such networks are called **deep networks**.

Autoencoder

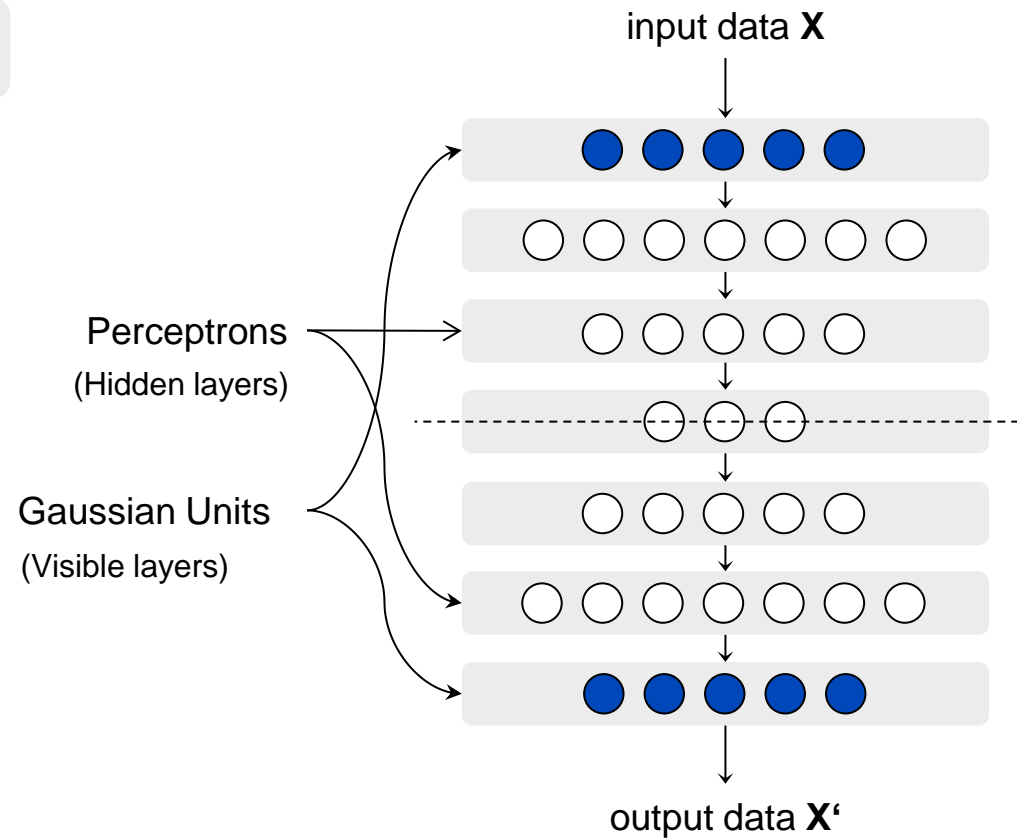
- Deep network



Such networks are called **deep networks**.

Autoencoder


- Deep network
- Symmetric topology



Autoencoders have a **symmetric topology** ...

Autoencoder

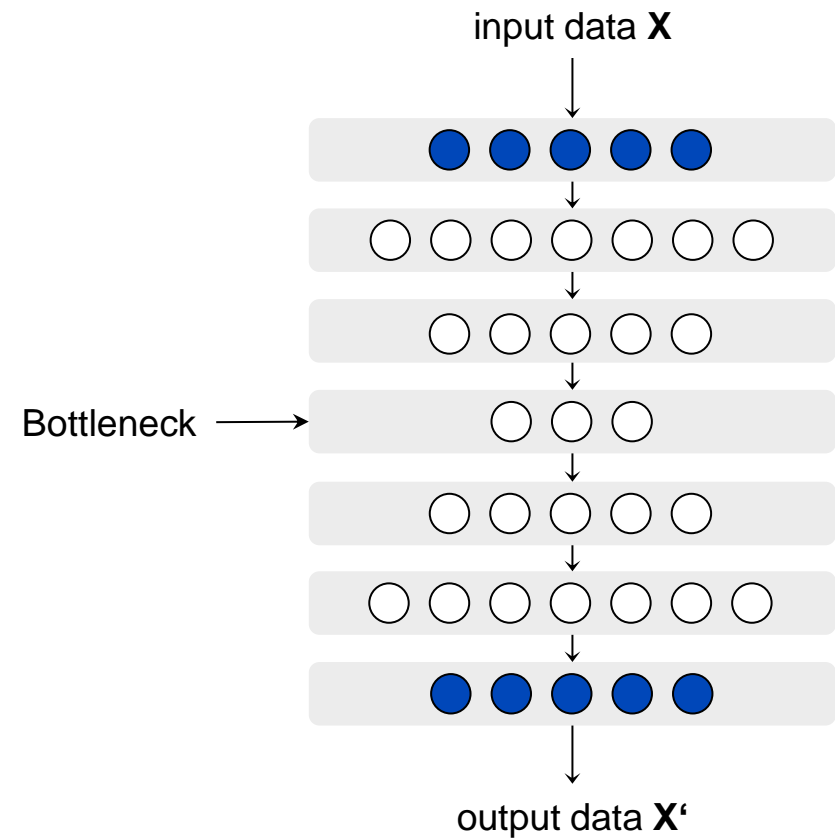
- Deep network
- Symmetric topology



... with an **odd number** of hidden layers.

Autoencoder


- Deep network
- Symmetric topology
- Information bottleneck



The small layer in the center works like an **information bottleneck**

Autoencoder


- Deep network
- Symmetric topology
- Information bottleneck



... that creates a **low dimensional code** for each sample in the input data.

Autoencoder

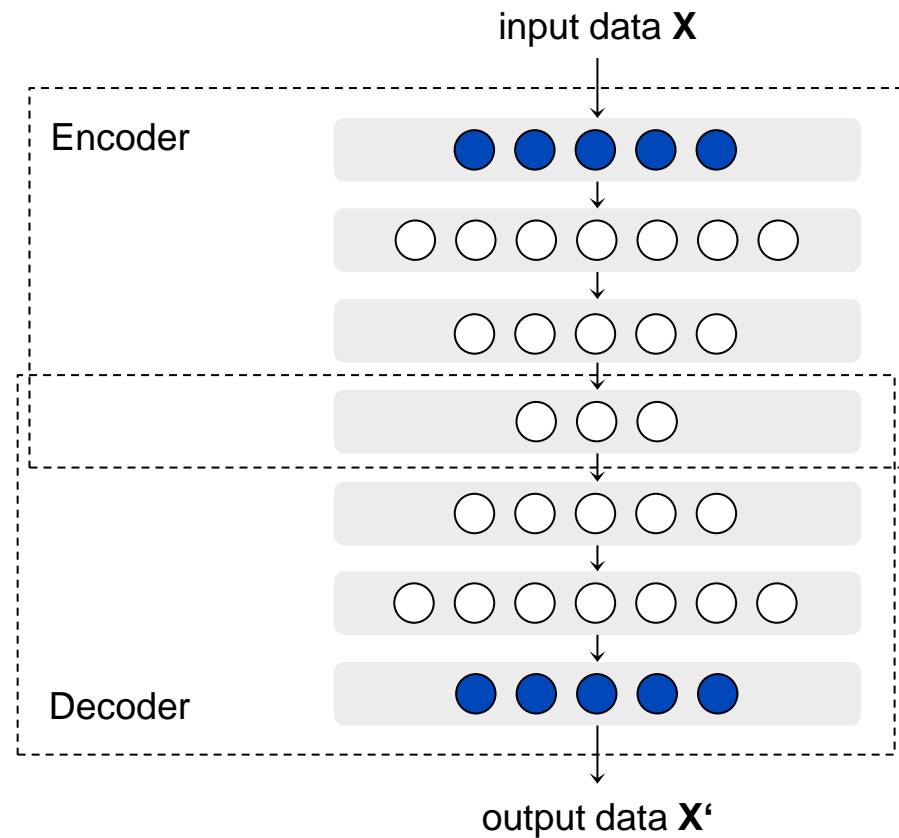
- Deep network
- Symmetric topology
- Information bottleneck
- Encoder



The upper stack does the **encoding** ...

Autoencoder

- Deep network
- Symmetric topology
- Information bottleneck
- Encoder
- Decoder



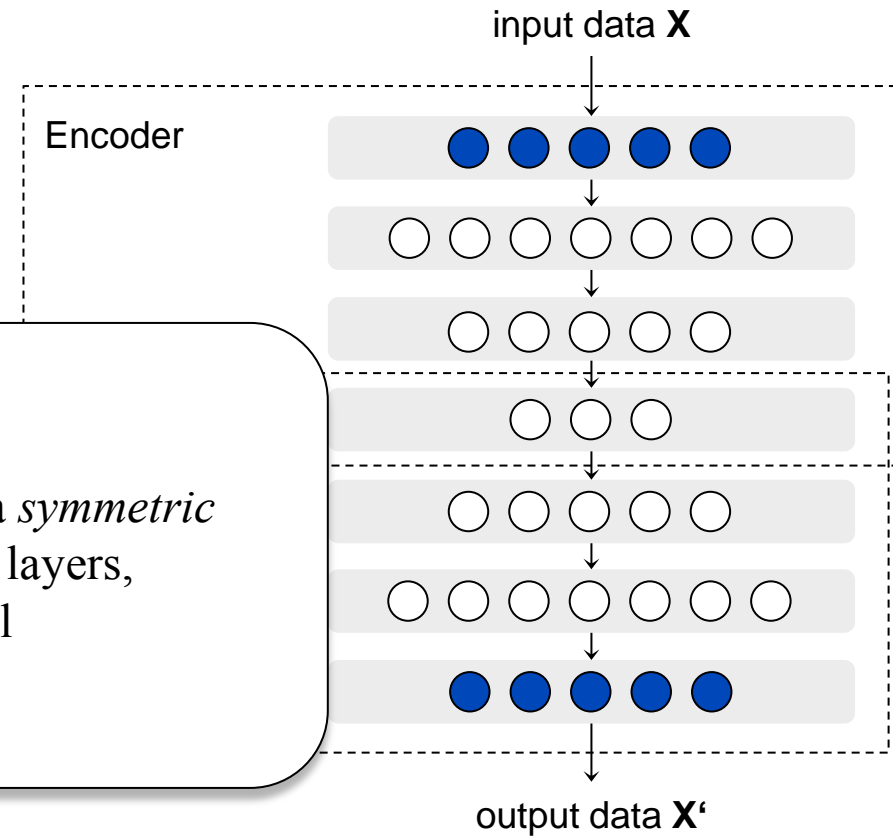
... and the lower stack does the **decoding**.

Autoencoder

- Deep network
- Symmetric topology
- Information bottleneck

Definition (*autoencoder*)

Autoencoders are *deep networks* with a *symmetric topology* and an odd number of hidden layers, containing a *encoder*, a low dimensional representation and a *decoder*.




... and the lower stack does the **decoding**.

Autoencoder

Problem: dimensionality of data

Idea:

1. Train autoencoder to minimize the distance between input \mathbf{X} and output \mathbf{X}'
2. Encode \mathbf{X} to low dimensional code \mathbf{Y}
3. Decode low dimensional code \mathbf{Y} to output \mathbf{X}'
4. Output \mathbf{X}' is low dimensional




Autoencoders can be used to **reduce the dimension of data** ...

Autoencoder

Problem: dimensionality of data

Idea:

1. Train autoencoder to minimize the distance between input \mathbf{X} and output \mathbf{X}'
2. Encode \mathbf{X} to low dimensional code \mathbf{Y}
3. Decode low dimensional code \mathbf{Y} to output \mathbf{X}'
4. Output \mathbf{X}' is low dimensional




... if we can train them!

Autoencoder

Training

Backpropagation



In feedforward ANNs **backpropagation** is a good approach.

Autoencoder


Training

Backpropagation

- (1) The distance (error) between current output X' and wanted output Y is computed. This gives a error function

$$X' = F(X)$$
$$\text{error} = \sqrt{X'^2 - Y}$$

output data X'



In feedforward ANNs **backpropagation** is a good approach.

Autoencoder



Training

input data \mathbf{X}

Backpropagation

- (1) The distance (error) between current output \mathbf{X}' and wanted output \mathbf{Y} is computed. This gives a error function

Example (*linear neuronal unit with two inputs*)



Autoencoder

Training

B

Backpropagation

- (1) The distance (error) between current output \mathbf{X}' and wanted output \mathbf{Y} is computed. This gives a error function
- (2) By calculating $-\nabla error$ we get a vector that shows in a direction which decreases the error
- (3) We update the parameters to decrease the error

input data \mathbf{X}




output data \mathbf{X}'

In feedforward ANNs **backpropagation** is a good approach.

Autoencoder


Training

input data \mathbf{X}



Backpropagation

- (1) The distance (error) between current output \mathbf{X}' and wanted output \mathbf{Y} is computed. This gives a error function
- (2) By calculating $-\nabla error$ we get a vector that shows in a direction which decreases the error
- (3) We update the parameters to decrease the error
- (4) We repeat that




Autoencoder

Training

Backpropagation

Problem: Deep Network



... the problem are the multiple hidden layers!


Autoencoder

Training

Backpropagation

Problem: Deep Network

- Very slow training



Backpropagation is known to be slow far away from the output layer ...


Autoencoder

Training

Backpropagation

Problem: Deep Network

- Very slow training
- Maybe bad solution



... and can converge to poor **local minima**.

Autoencoder


Training

Backpropagation

Problem: Deep Network

- Very slow training
- Maybe bad solution

Idea: Initialize close to a good solution



The task is to **initialize the parameters** close to a good solution!

Autoencoder

Training


Backpropagation

Problem: Deep Network

- Very slow training
- Maybe bad solution

Idea: Initialize close to a good solution

- Pretraining



Therefore the training of autoencoders has a **pretraining** phase ...

Autoencoder

Training


Backpropagation

Problem: Deep Network

- Very slow training
- Maybe bad solution

Idea: Initialize close to a good solution

- Pretraining
- Restricted Boltzmann Machines



... which uses **Restricted Boltzmann Machines (RBMs)**

Autoencoder

input data \mathbf{X}

Restricted Boltzmann Machine

- RBMs are **Markov Random Fields**

Bas

Pr

-
-

Id

-
-

Autoencoder

input data X

Restricted Boltzmann Machine


- RBMs are **Markov Random Fields**

Markov Random Field

- Every unit influences every neighbor
- The coupling is undirected

Motivation (Ising Model)

A set of magnetic dipoles (*spins*) is arranged in a graph (lattice) where neighbors are coupled with a given strength




Autoencoder

input data \mathbf{X}

Restricted Boltzmann Machine

- RBMs are **Markov Random Fields**
- Bipartite topology: **visible** (v), **hidden** (h)
- Use local **energy** to calculate the probabilities of values

Training:
contrastive divergency
(Gibbs Sampling)

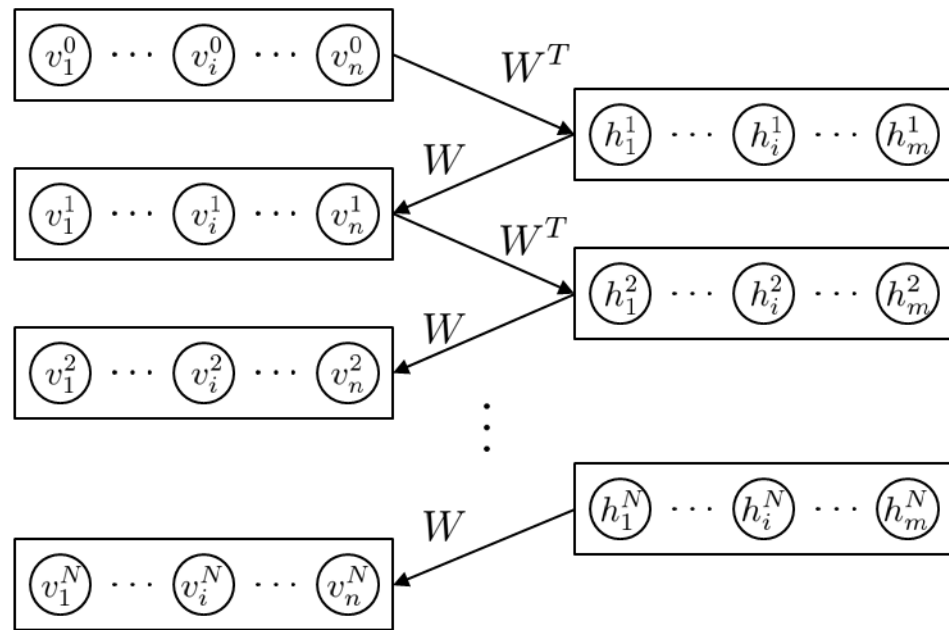


Autoencoder

input data \mathbf{X}

Restricted Boltzmann Machine

Gibbs Sampling



Autoencoder

Training

Top

V := set of visible units

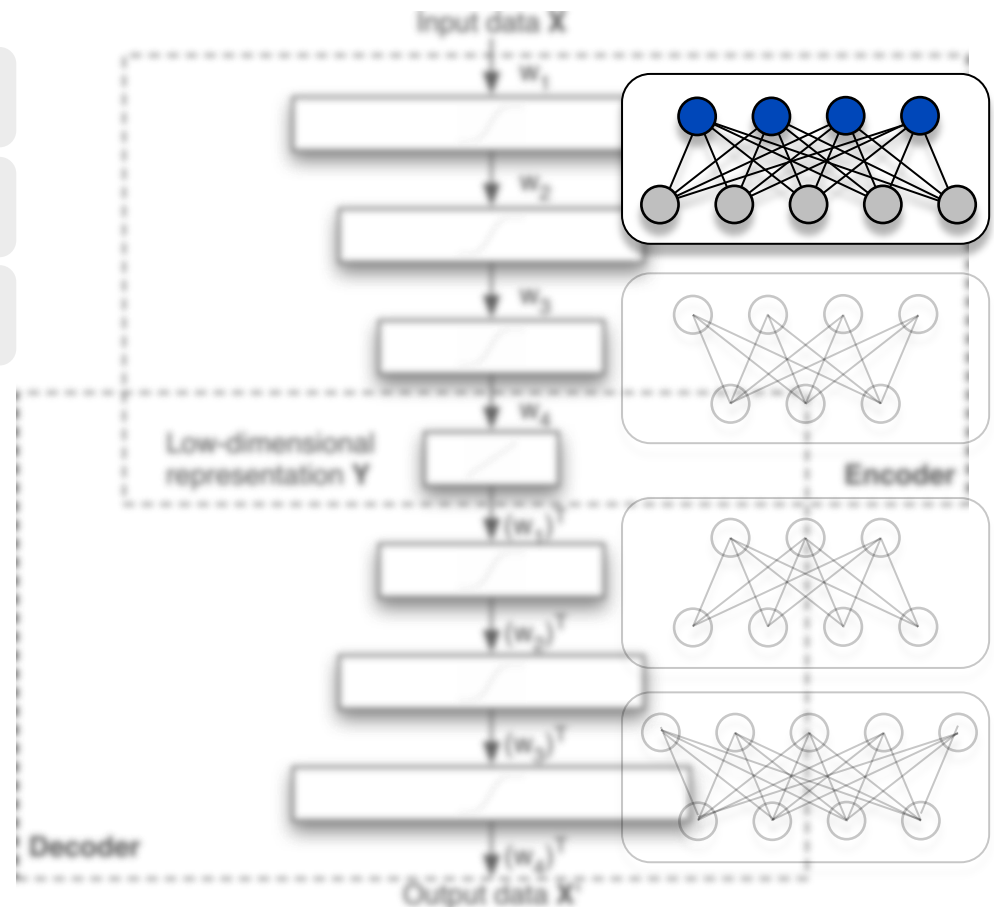
x_v := value of unit v , $\forall v \in V$

$x_v \in \mathbf{R}$, $\forall v \in V$

H := set of hidden units

x_h := value of unit h , $\forall h \in H$

$x_h \in \{0, 1\}$, $\forall h \in H$



The top layer RBM transforms **real value data** into binary codes.

Autoencoder

Training




Top

$$x_v \sim N\left(b_v + \sum_h w_{vh} x_h, \sigma_v\right)$$

σ_v := std. dev. of unit v

b_v := bias of unit v

w_{vh} := weight of edge (v, h)



Therefore visible units are modeled with **gaussians** to encode **data** ...

Autoencoder

Training




Top

$$x_h \sim \text{sigm} \left(b_h + \sum_v w_{vh} \frac{x_v}{\sigma_v} \right)$$

σ_v := std. dev. of unit v

b_h := bias of unit h


w_{vh} := weight of edge (v, h)



... and many hidden units with **simoids** to encode **dependencies**

Autoencoder

Training




Top

Local Energy

$$E_v := - \sum_h \underset{\text{green}}{w_{vh}} \underset{\text{red}}{\frac{x_v}{\sigma_v}} x_h + \frac{(x_v - \underset{\text{blue}}{b_v})^2}{2 \underset{\text{red}}{\sigma_v}^2}$$

$$E_h := - \sum_v \underset{\text{green}}{w_{vh}} \underset{\text{red}}{\frac{x_v}{\sigma_v}} x_h + x_h \underset{\text{blue}}{b_h}$$



The **objective function** is the sum of the local energies.

Autoencoder

Training

Reduction

V := set of visible units


x_v := value of unit v , $\forall v \in V$

$x_v \in \{0, 1\}$, $\forall v \in V$

H := set of hidden units

x_h := value of unit h , $\forall h \in H$

$x_h \in \{0, 1\}$, $\forall h \in H$



The next RBM layer **maps** the dependency encoding...

Autoencoder

Training




Reduction

$$x_v \sim \text{sigm} \left(b_v + \sum_h w_{vh} x_h \right)$$

b_v := bias of unit v

w_{vh} := weight of edge (v, h)



... from the upper layer ...

Autoencoder

Training




Reduction

$$x_h \sim \text{sigm} \left(b_h + \sum_v w_{vh} x_v \right)$$

b_h := bias of unit h

w_{vh} := weight of edge (v, h)



... to a smaller number of **simoids** ...

Autoencoder

Training


Reduction

$$\hat{f}_i = \int_{\mathcal{H}} f_i(\mathbf{h}) d\mathbf{h}$$

Local Energy

$$E_v := - \sum_h \mathbf{w}_{vh} x_v x_h + x_h \mathbf{b}_h$$

$$E_h := - \sum_v \mathbf{w}_{vh} x_v x_h + x_v \mathbf{b}_v$$




... which can be trained faster than the top layer

Autoencoder

Training

Unrolling




The **symmetric topology** allows us to skip further training.

Autoencoder

Training

Unrolling




The **symmetric topology** allows us to skip further training.

Autoencoder

Training

- **Pretraining**
Top RBM (GRBM)
Reduction RBMs
Unrolling
- **Finetuning**
Backpropagation

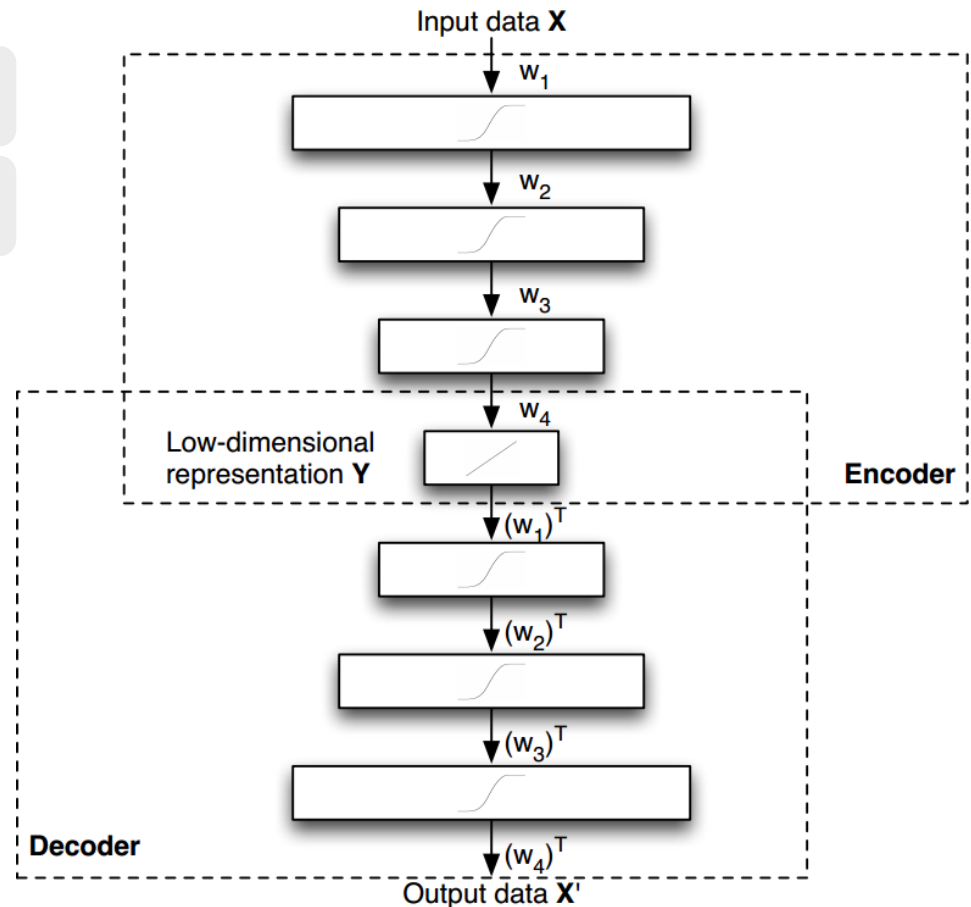


After pretraining **backpropagation** usually finds **good solutions**


Autoencoder

Training


- **Complexity:** $O(inw)$
i: number of iterations
n: number of nodes
w: number of weights
- **Memory Complexity:** $O(w)$



The **algorithmic complexity** of RBM training depends on the network size



Autoencoders




Biological Model



Validation & Implementation

Network Modeling


Restricted Boltzmann Machines (RBM)



How to model the topological structure?

Network Modeling


Restricted Boltzmann Machines (RBM)



We define S and E as **visible data Layer** ...

Network Modeling


Restricted Boltzmann Machines (RBM)



We identify S and E with the **visible layer** ...

Network Modeling


Restricted Boltzmann Machines (RBM)




... and the TFs with the **hidden layer** in a RBM

Network Modeling


Restricted Boltzmann Machines (RBM)



The training of the RBM gives us a model



Autoencoder



Biological Model

```
class AutoEncoder:
    def __init__(self,
                 num_hidden = 10,
                 num_visible = 100,
                 learning_rate = 0.1):
        # Initialize a weight matrix
        # a Gaussian distribution
        self.weights = 0.1 * np.random.randn(num_hidden, num_visible)
        # Insert weights
```

Implementation & Results

Validation of the results

- Needs information about the true regulation
- Needs information about the descriptive power of the data

Validation of the results

- Needs information about the true regulation
- Needs information about the descriptive power of the data

Without this information validation can only be done,
using **artificial datasets!**

Artificial datasets

We simulate data in three steps:

Artificial datasets

We simulate data in three steps

Step 1

Choose number of Genes ($E+S$) and create random bimodal distributed data

Artificial datasets

We simulate data in three steps

Step 1

Choose number of Genes ($E+S$) and create random bimodal distributed data

Step 2

Manipulate data in a fixed order

Artificial datasets

We simulate data in three steps

Step 1

Choose number of Genes ($E+S$) and create random bimodal distributed data

Step 2

Manipulate data in a fixed order

Step 3

Add noise to manipulated data
and normalize data


Simulation

Step 1

Number of visible nodes 8 (4E, 4S)

Create random data:

Random $\{-1, +1\} + N(0, \sigma = 0.5)$



Simulation

Step 2

Manipulate data

$$e_1 = 0.25s_1 + 0.25s_2 + 0.25s_3 + 0.25s_4$$

$$e_2 = 0.5s_1 + 0.5 \text{ Noise}$$


$$e_3 = 0.5s_1 + 0.5 \text{ Noise}$$

$$e_4 = 0.5s_1 + 0.5 \text{ Noise}$$

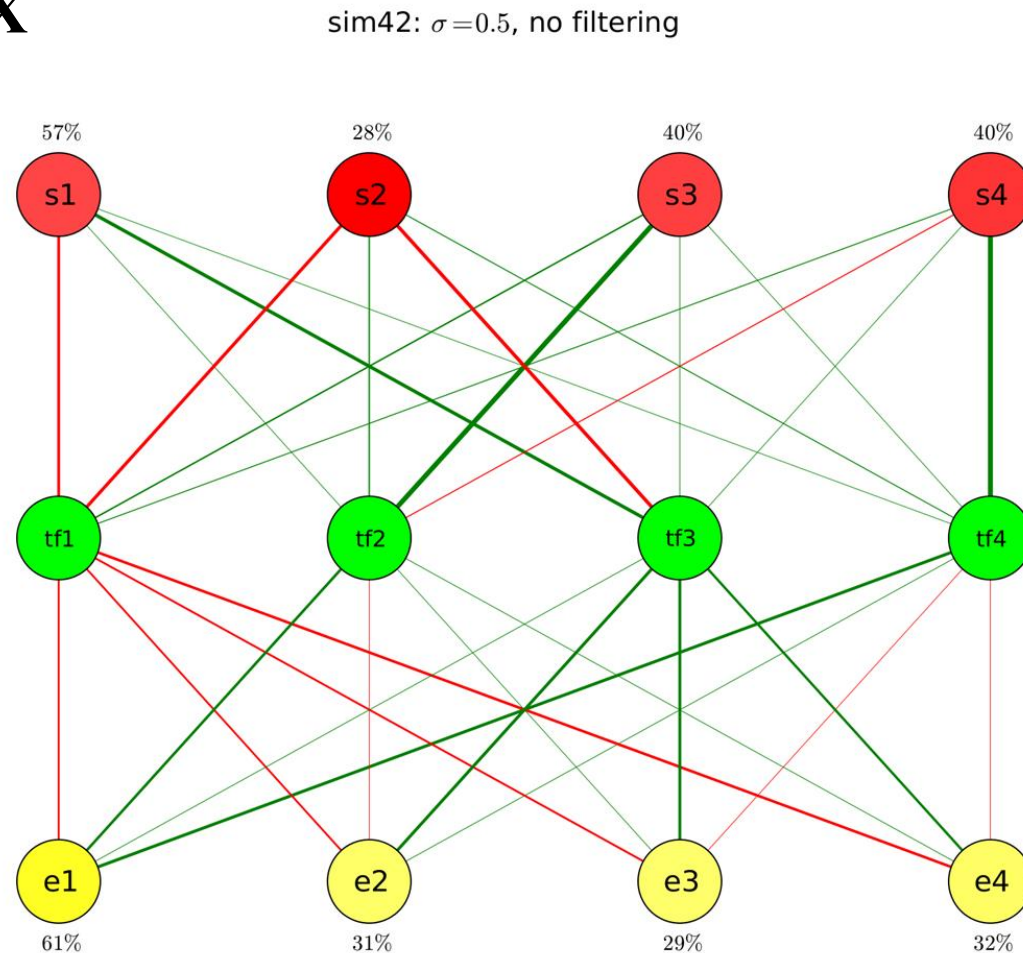
Simulation

Step 3

Add noise: $N(0, \sigma = 0.5)$



We analyse the data **X**
with an RBM



Average performance: 40.3%

We train an autoencoder with 9 hidden layers and 165 nodes:

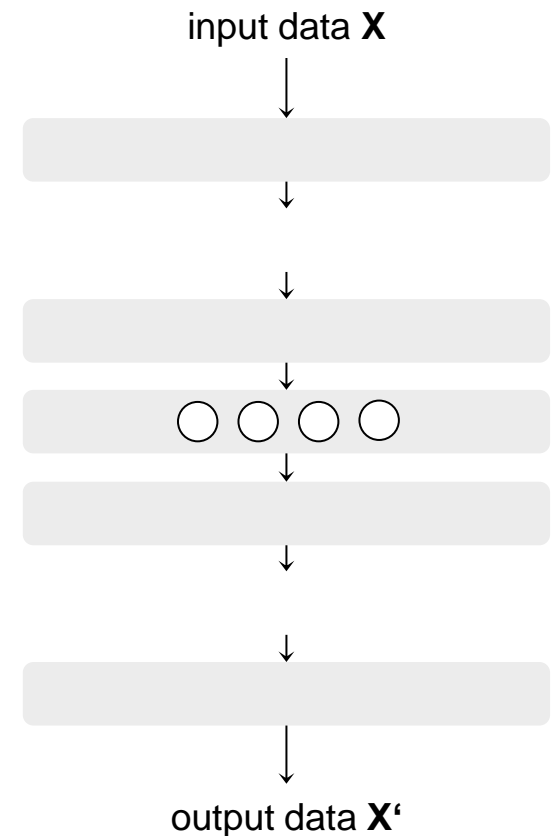
Layer 1 & 9: 32 hidden units

Layer 2 & 8: 24 hidden units



Layer 3 & 7: 16 hidden units

Layer 4 & 6: 8 hidden units

Layer 5: 5 hidden units




We transform the data from \mathbf{X} to \mathbf{X}'
And reduce the dimensionality



We analyse the
transformed data \mathbf{X}'
with an RBM


sim42: $\sigma=0.5$, (32,24,16,8,4) filtering



Average performance: 69.5%


Lets compare the models

sim42: $\sigma=0.5$, no filtering



Average performance: 40.3%


sim42: $\sigma=0.5$, (32,24,16,8,4) filtering



Average performance: 69.5%


Another Example with more nodes and larger autoencoder

sim40: $\sigma=0.5$, no filtering



Average performance: 50.6%

sim40: $\sigma=0.5$, (64,48,32,16,8) filtering



Average performance: 100.0%

Conclusion

- Autoencoders can improve modeling significantly by **reducing the dimensionality of data**
- Autoencoders **preserve complex structures** in their multilayer perceptron network. Analysing those networks (for example with knockout tests) could give more structural information
- The drawback are **high computational costs**
Since the field of deep learning is getting more popular (Face recognition / Voice recognition, Image transformation). Many new improvements in facing the computational costs have been made.

Acknowledgement

eilsLABS

Prof. Dr. Rainer König

Prof. Dr. Roland Eils

Network Modeling Group

